

# Make RPC Faster than Fast

A detailed look at the performance optimization journey of Kitex to improve RPC speed.

by Zhuowei Wang  
CloudWeGo Kitex&Netpoll Maintainer



# Content

# 目录

**01** The Costs of Microservices

**02** Reducing CPU Cost

**03** Reducing Latency Cost

**04** Benchmark






# Part 01

## The Costs of Microservices



## Monolithic Architecture

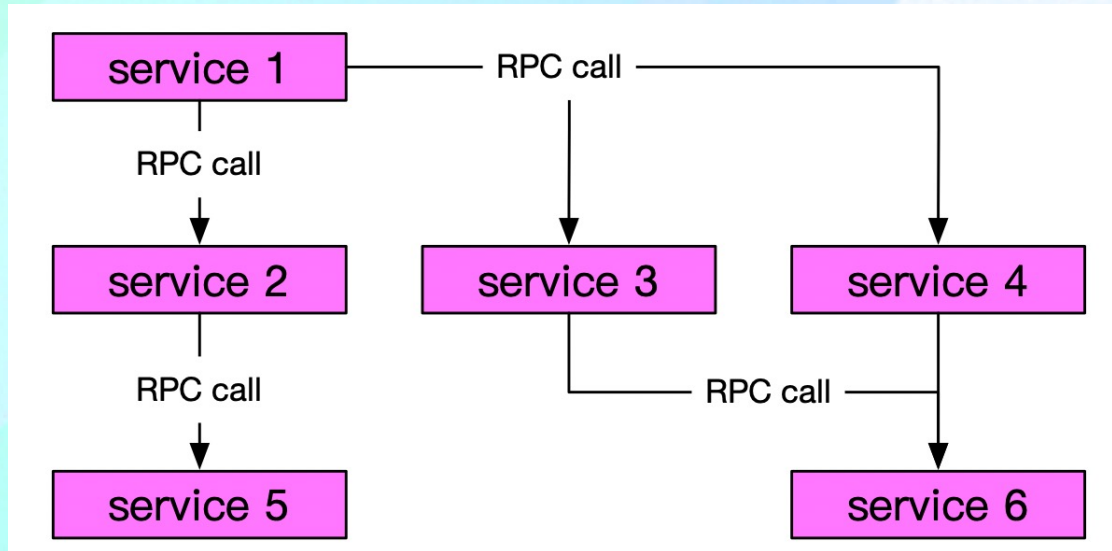


Function Call Cost: ~10 ns

But: Significant Maintainability and Stability Costs

- multiple teams develop in a same codebase
- If 1 module crash, the whole service process crash

## Microservices Architecture



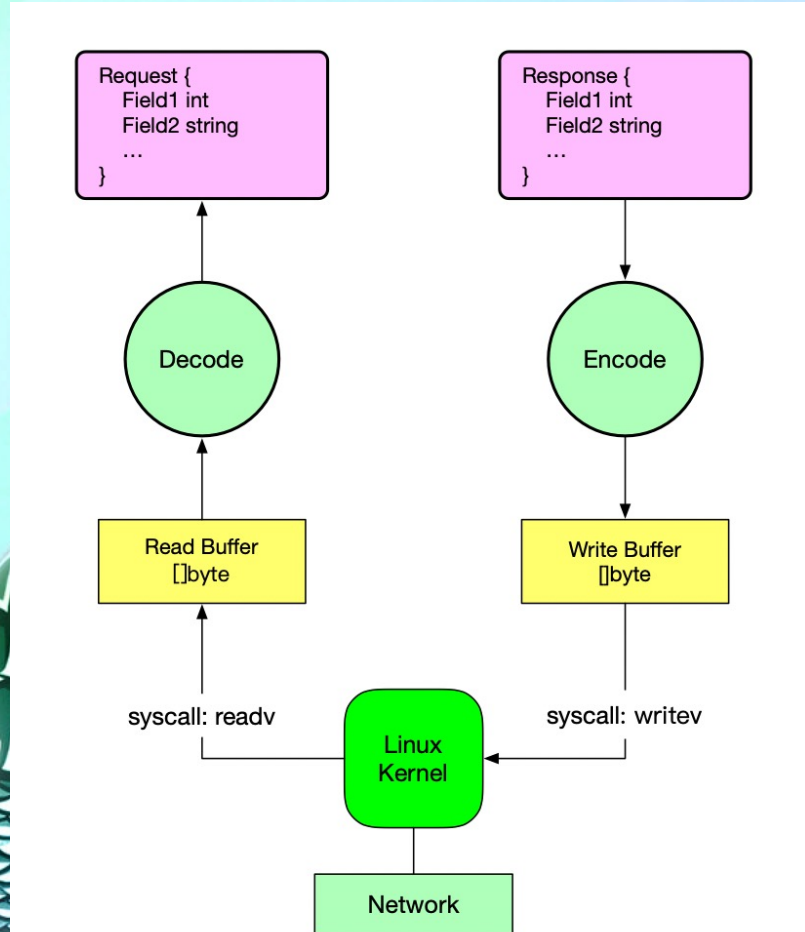
Every service have more Maintainability and Stability

- Different teams are responsible for different services and don't need to worry about other services
- If an RPC call failed, can retry to call another downstream instance

But: RPC Call Cost: 1~100 ms



## What constitutes the Cost of RPC



### Costs:

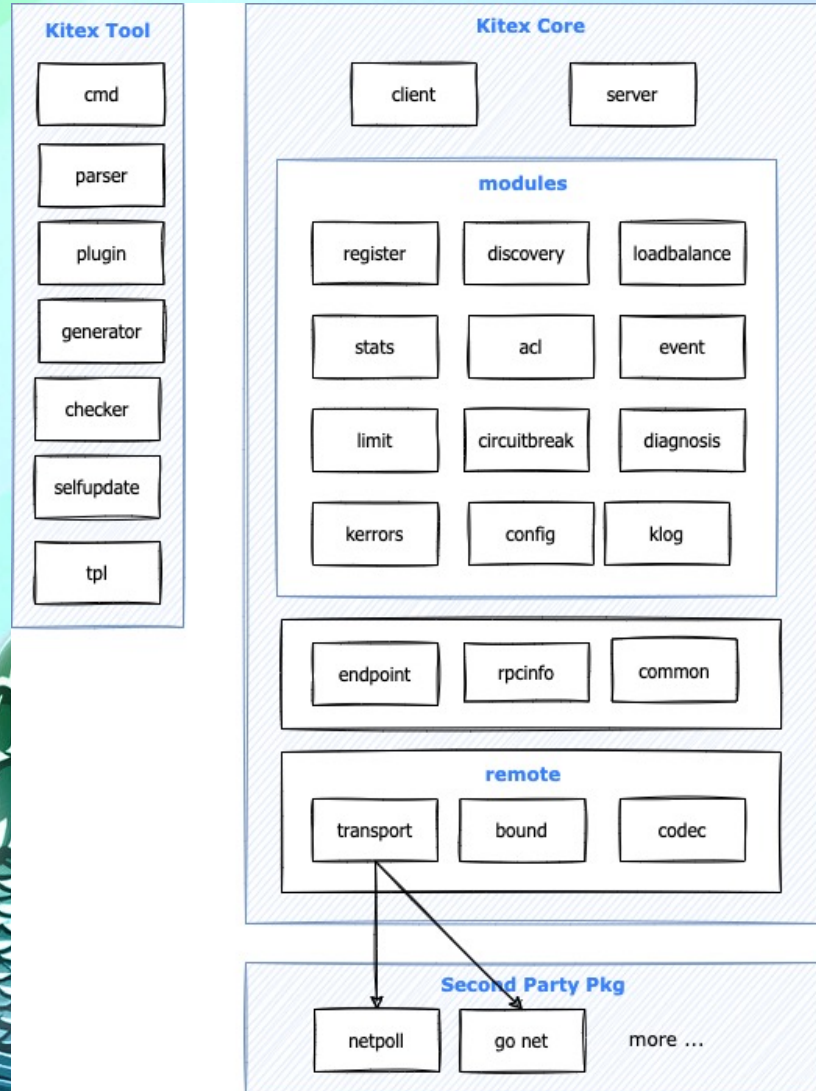
- **CPU Cost**

- Codec (Decode/Encode)
- IO syscall (read/write)
- .....

- **Latency Cost**

- Kernel thread scheduling
- Network transmission
- .....

## What is KiteX?



**KiteX** is a high-performance and strong-extensibility Golang RPC framework that helps developers to build microservices.

# Part 02

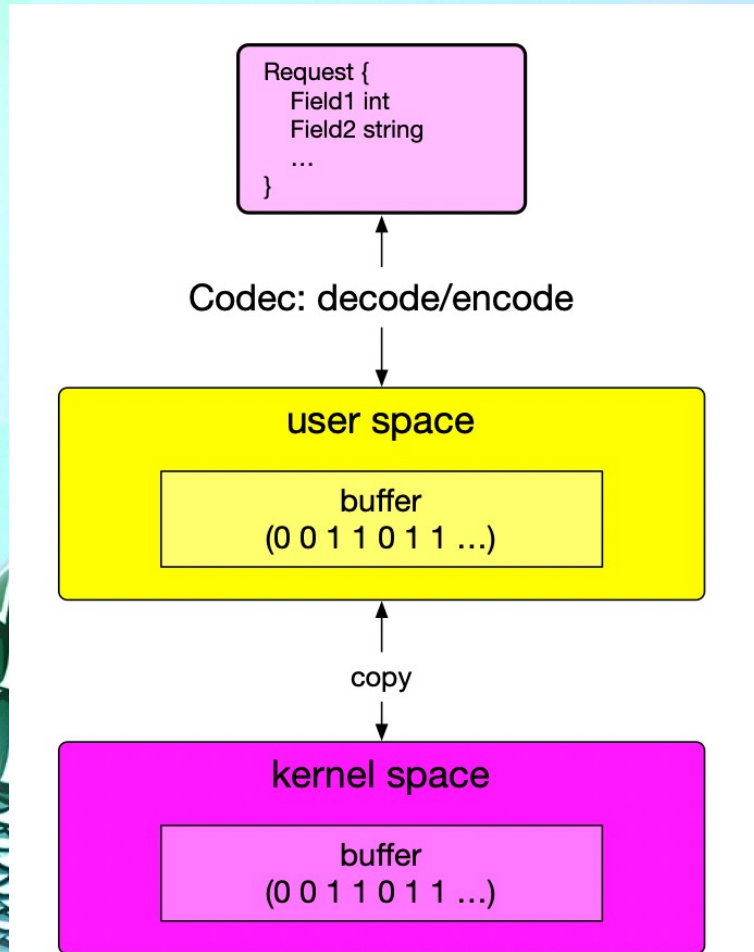
## Reducing CPU Cost

The Optimizations of Kitex Codec





## What is codec, and why we need it?



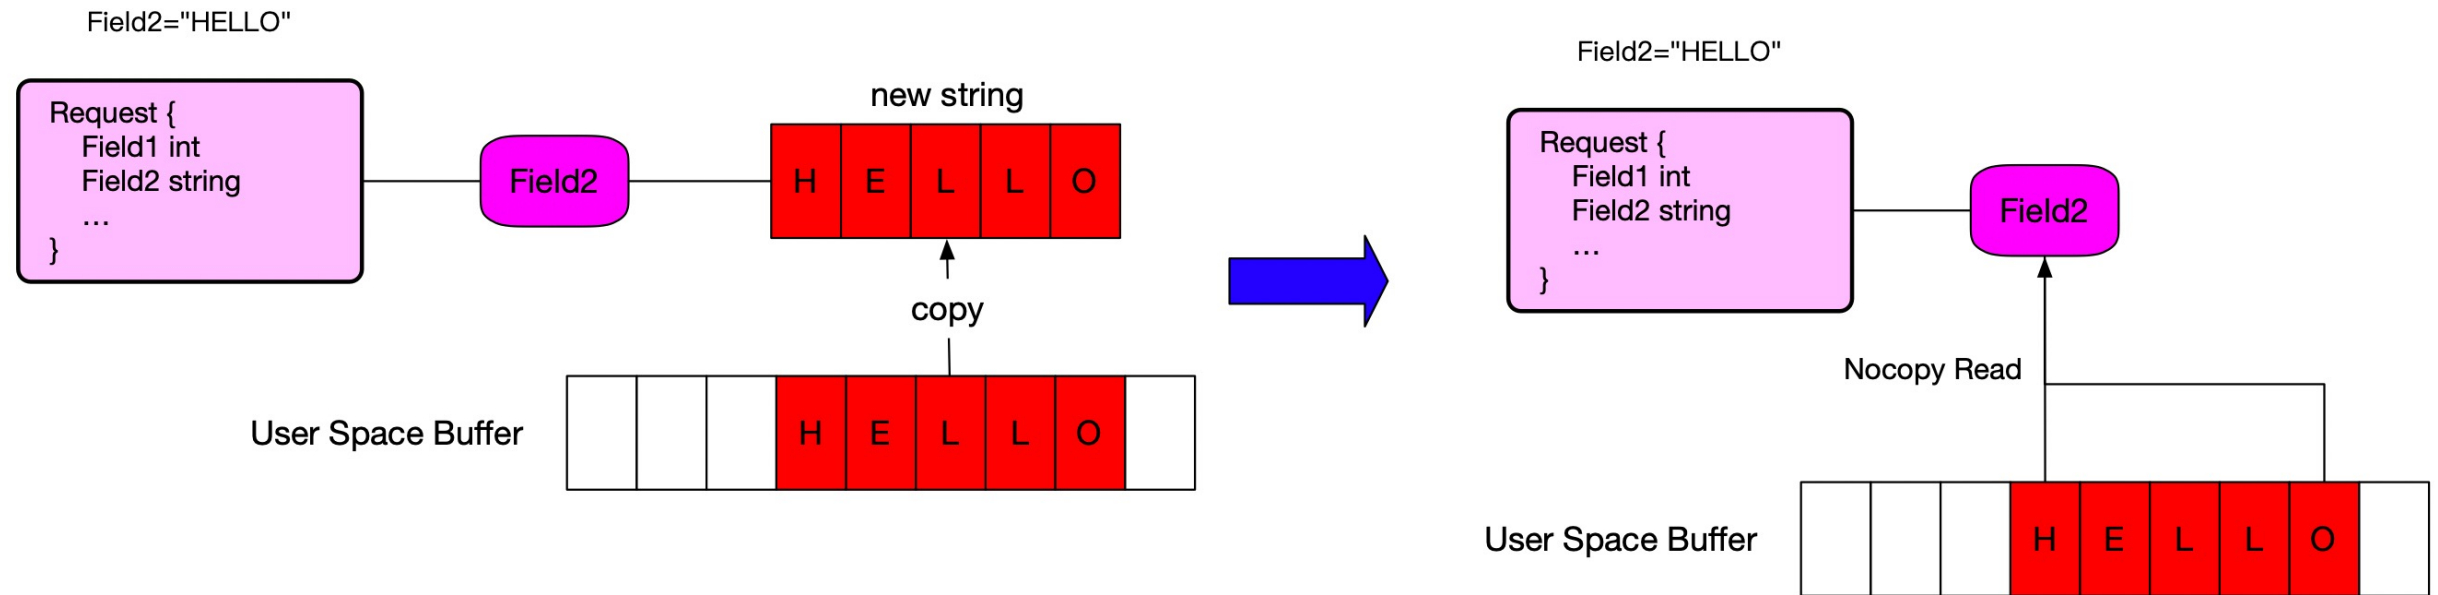
Codec helps to translate data view between:

Human View  $\Leftrightarrow$  Machine View

The codec is highly CPU-intensive, with our online machines consuming over 20% of CPU usage for codec.



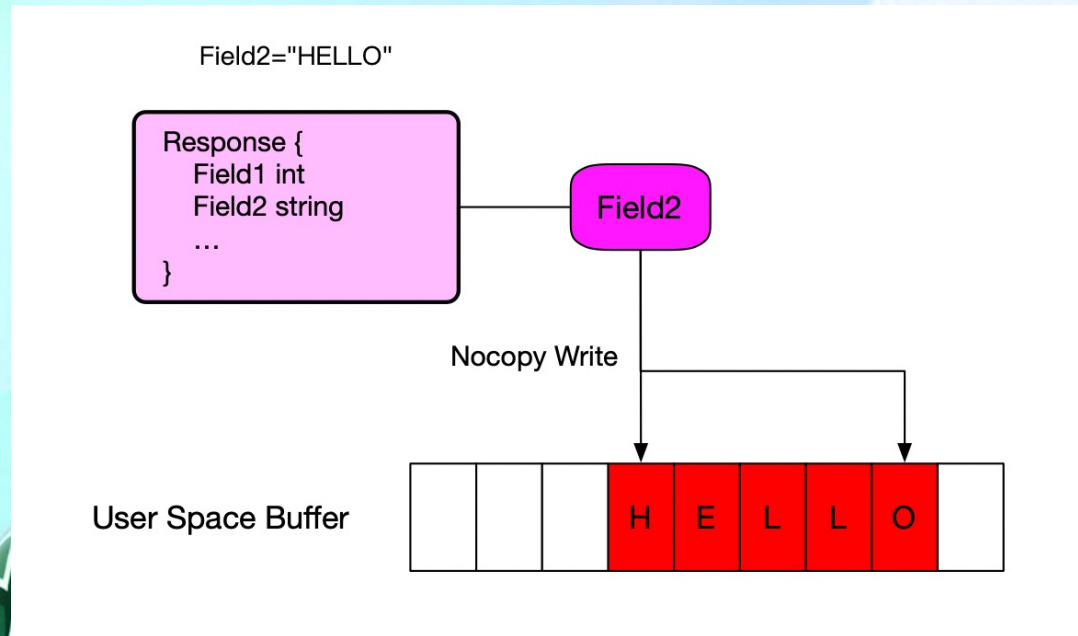
## Kitex Codec Optimization: NoCopy Read



Nocopy read for string/bytes type fields

``request.Field2`` uses the user space buffer memory directly, without any malloc

## Kitex Codec Optimization: NoCopy Write

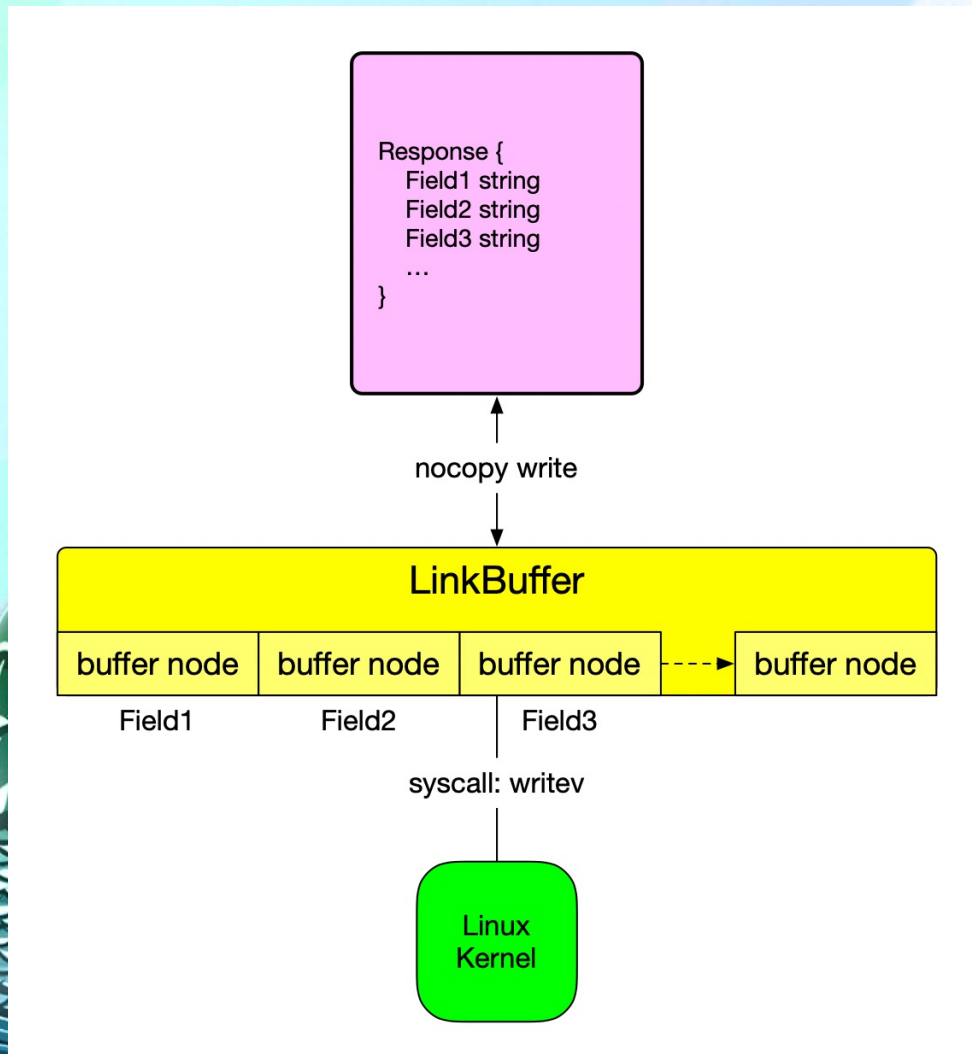


Nocopy write for string/bytes type fields

The user space buffer uses the memory of `request.Field2` directly, without any malloc



## Kitex Codec Optimization: LinkBuffer




**LinkBuffer:** The Key to No-Copy Read/Write

**writev** syscall: Write multiple buffer nodes with a single system call

## Kitex Codec Optimization: Reduce Object Malloc

```
Request {  
  container []*Object  
  ...  
}
```

A: container := make([]\*Object, N)



B: container := make([]Object, N)



code A: N Malloc calls and N pointers

code B: 1 Malloc calls and 1 big pointer

### Why code B is better?

1. GC only need scan the big pointer once
2. The contiguous memory address make CPU cache happy

# Part 03

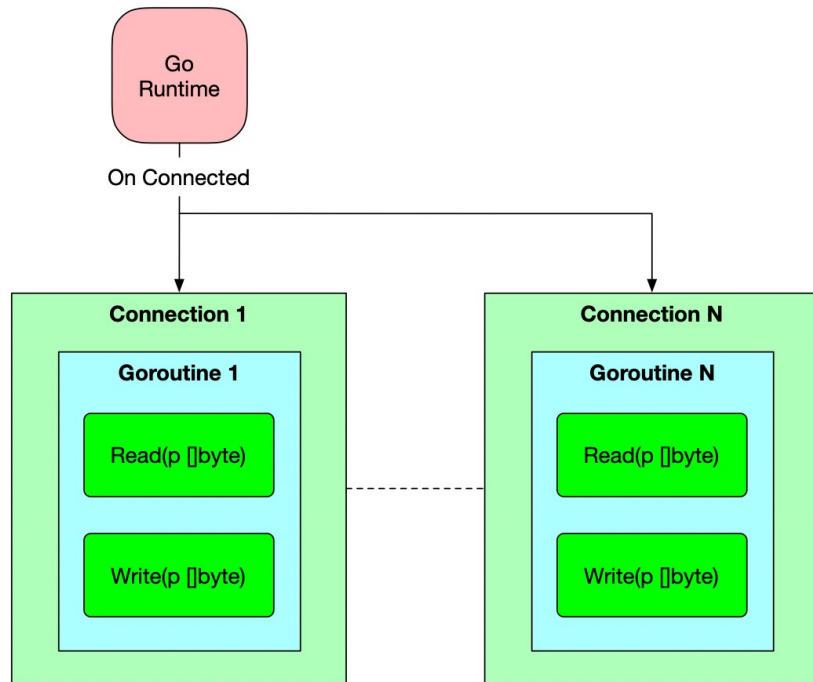
## Reducing Latency Cost

The Optimizations of Kitex network layer



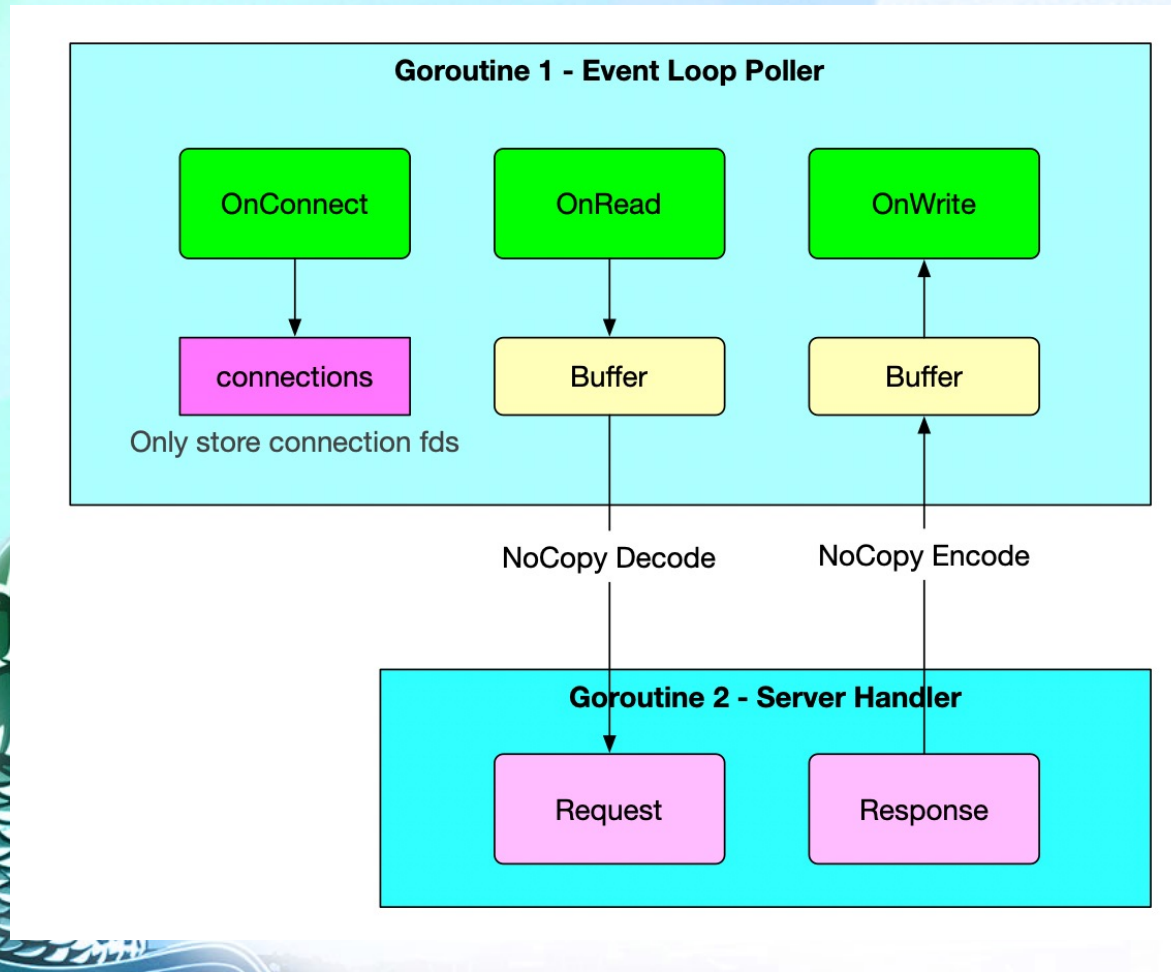


## The Issues with Go's Native Net Library



- N connections create N goroutines, leading to high goroutine scheduling latency.
- The native Read/Write API does not support NoCopy Read and NoCopy Write abilities.

## cloudwego/netpoll: An alternative to the Go net library



### Higher Throughput

- Million Connections with zero Goroutines
- Support No Copy Codec

### Lower Latency

- Independently from Go runtime scheduler
- Read/Write data within the same goroutine to enhance cache efficiency

# Part 04

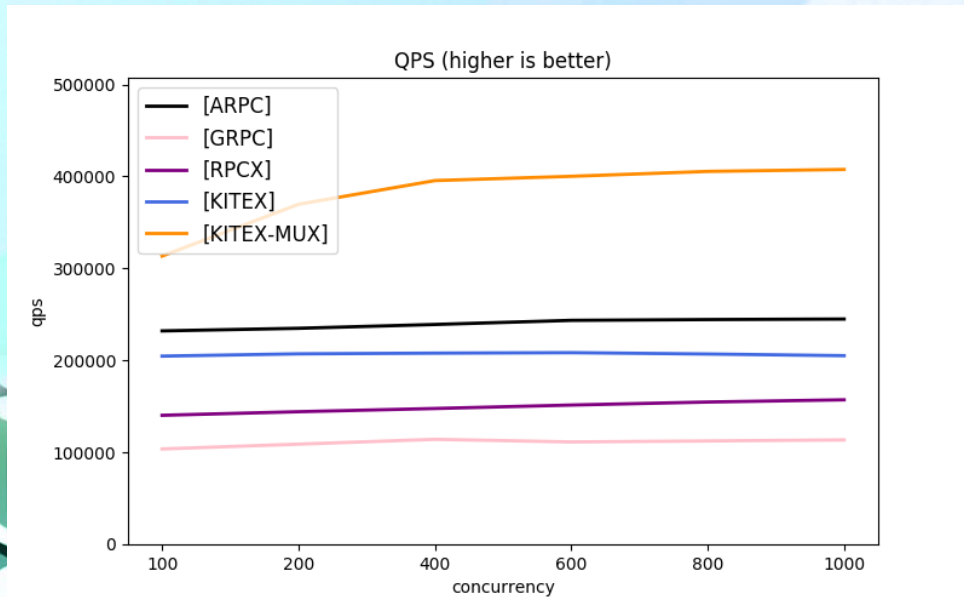
## Benchmark






## The Benchmark comparison of Kitex with other RPC frameworks.

### Highest Throughput



### lowest Latency



KITEX: kitex in long connection mode  
KITEX\_MUX: kitex in multiplexing mode

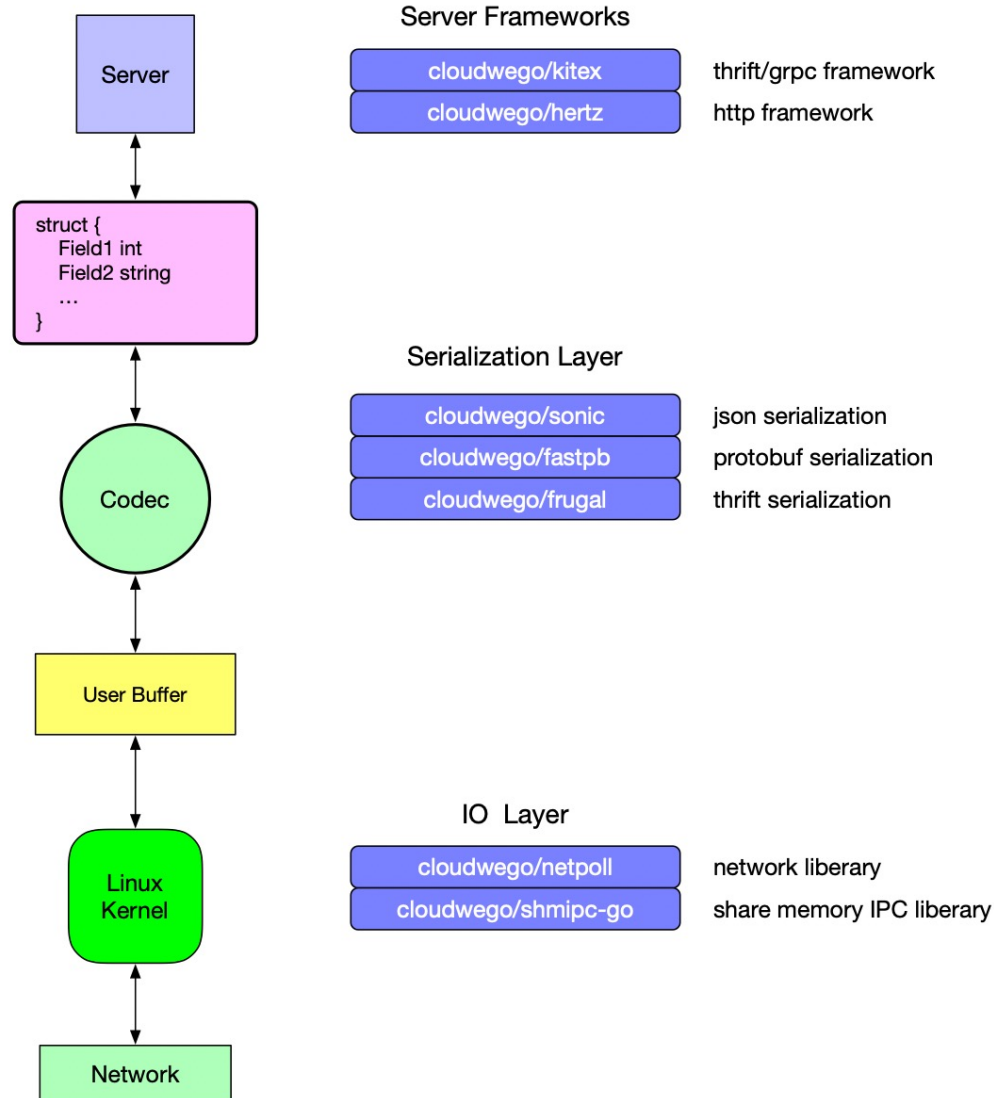
data source: <https://github.com/cloudwego/kitex-benchmark>

What's more?



# Cloud Native Microservices Meetup

## CloudWeGo Family



### Server Frameworks

- <https://github.com/cloudwego/kitex>
- <https://github.com/cloudwego/hertz>

### Serialization Library:

- <https://github.com/cloudwego/sonic>
- <https://github.com/cloudwego/frugal>
- <https://github.com/cloudwego/fastpb>

### IO Library:

- <https://github.com/cloudwego/netpoll>
- <https://github.com/cloudwego/shmipc-go>



**Thanks.**

